



CarbonSage

Track. Reduce. Sustain.

By Team elasticSearch

For **Hacknosis: Sustainable MedTech**




PROBLEM STATEMENT

The carbon footprint of our gadgets, the internet, and the essential systems supporting them accounts for 3.7% of global greenhouse emissions – on par with the airline industry. These emissions are predicted to double by 2025 as more devices hit the market, including smart medical devices.

Key Challenges:

- Increasing carbon emissions from technology usage.
- Need for effective monitoring and reduction strategies.
- Lack of awareness and tools for optimizing energy consumption.

OUR WORLD, NOW!



CO₂ Data of 2022 (tonnes)

CHINA	11,396,777,000.00
ASIA	10,375,572,000.00
EUROPE	5,105,307,700.00
USA	5,057,303,600.00
AFRICA	1,416,626,600.00
LATAM	1,085,392,400.00

REGION-WISE ANNUAL CO₂ EMISSIONS

Data source: Global Carbon Budget (2023) - [Learn more about this data](#)

SOLUTION OVERVIEW

Hackathon Objective: Develop a solution that can calculate, track, and reduce carbon emissions generated by devices or data usage, with auto-reporting using Opentext.

01

MONITOR

Track real-time energy usage and device performance.

02

FORECAST

Predict future energy consumption and identify inefficiencies.

03

REMEDY

Provide actionable insights and suggestions to reduce carbon emissions.

OUR 3 STEP APPROACH

KEY FEATURES




Anomaly Detection preventing Faults

For tracking excessive
CO2 emissions helping
prevent potential
Faults




Code Alternative

Analyze code block
carbon emissions and
suggest efficient code
alternatives using LLM.



CO2 Per Microservice

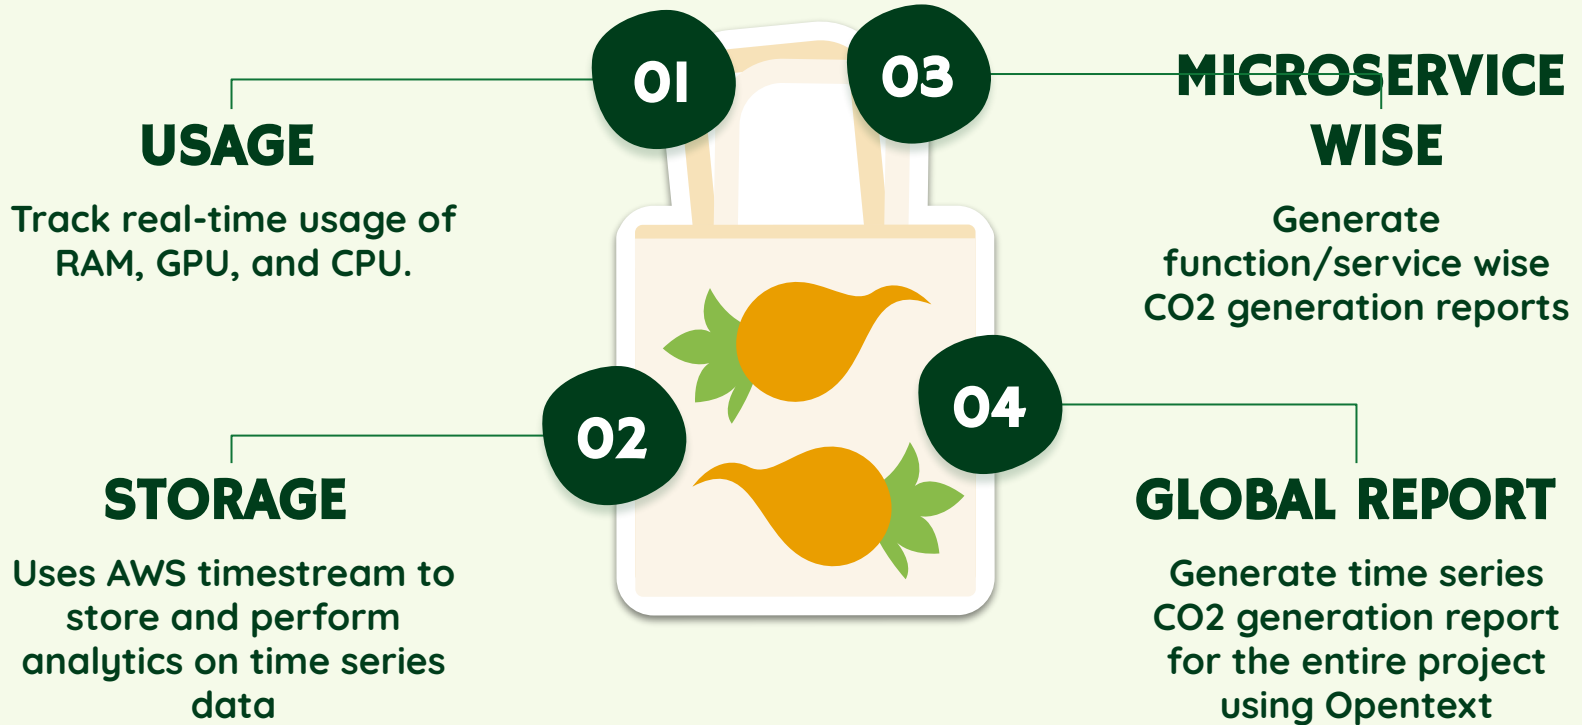
Calculating CO2
emission and energy
use on per
microservice basis



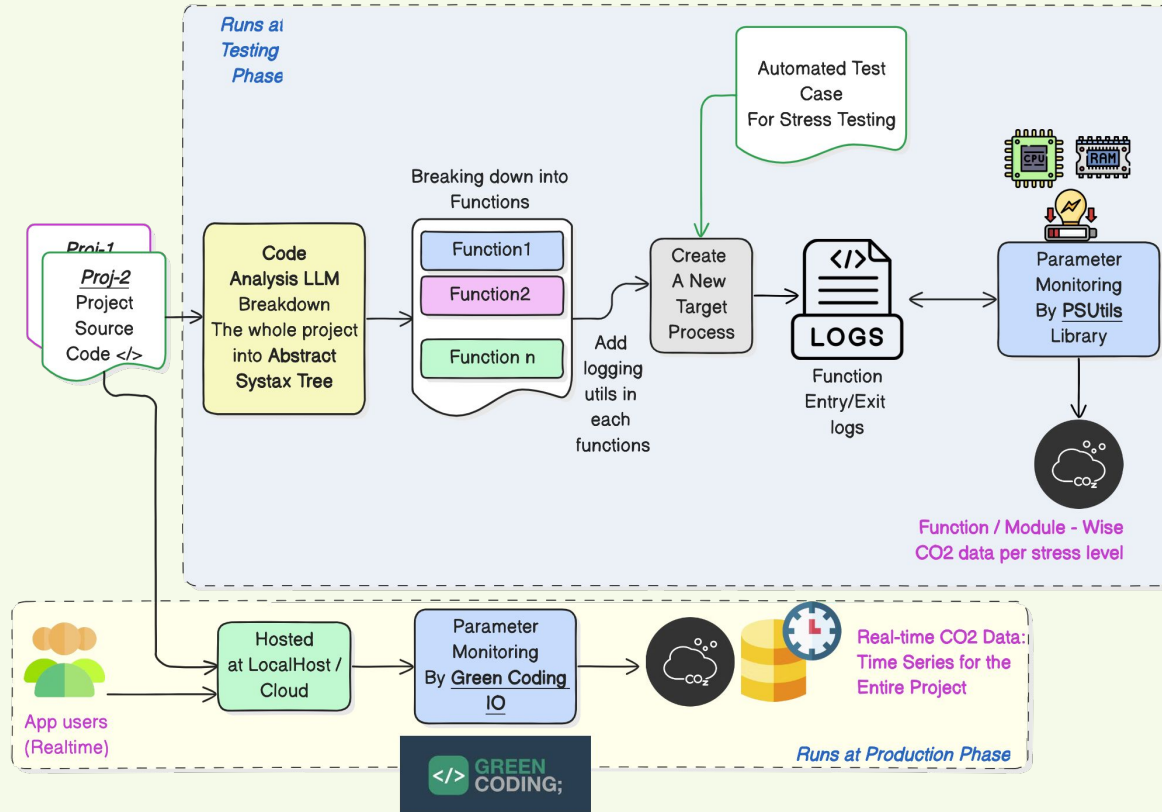
Usage Reporting

Generation of daily
usage reports on
service wise carbon
emissions

MONITORING



Data Collection and Storage (I/2)



Data Collection and Storage

(2/2)

We store the data mainly in two phases.

1. **Function/Module-Wise CO2 Data per Stress Level (Stored only at testing phase)**
 - a. We break the whole project into multiple microservices using code analysis LLMs such as CodeLlama and host them on different endpoints for stress testing.
 - b. We automatically generate test cases with respect to different stress levels.
 - c. We monitor different parameters like CPU usage, RAM usage, and power consumption using monitoring and tracking tools such as Green Coding IO.
2. **Real-time CO2 Data: Time Series for the Entire Project (Stored at production phase / realtime)**
 - a. We host the target project on cloud platforms such as AWS, GCP, or Localhost.
 - b. We monitor the system parameters with respect to the actual real-time user traffic.
 - c. We store the data in a time series format in AWS Timestream.



FORECAST

PREDICTION


Based on the CO2 time series data and service usage, predict growth of CO2 emissions

COST-BENEFIT ANALYSIS

Corroborate CO2 emission reductions with reduction in lines of code, code efficiency. Estimate effort for target reduction of CO2 emissions

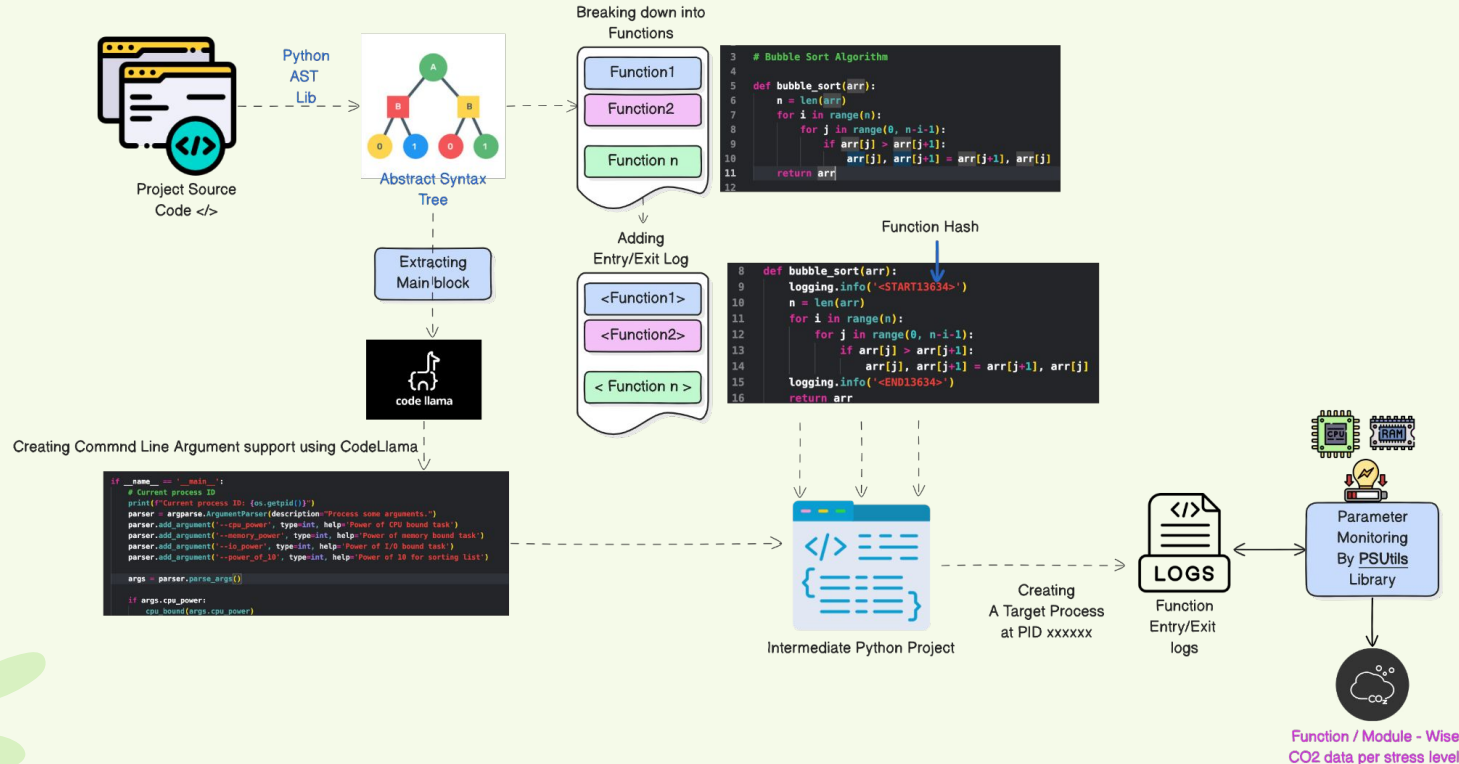
OPENTEXT ALERT NOTIFICATION

Based on abnormal or high CO2 emissions, notification is sent to project lead and writer of specific function/microservice



Code profiling for resource usage

Detailed Mechanism



Usage Analytics Generation

Parameter Monitoring Logs Generated By
PSUtils at 1 sec interval

1	cpu_percent	rss	vms	pfaults	pageins	time	iteration
2	63.9	12361728	34615349248	4265	73	2024-08-24 23:21:14	0
3	44.8	17432576	34618540032	5560	147	2024-08-24 23:21:15	1
4	45.5	23371776	34628603904	7319	682	2024-08-24 23:21:16	2
5	20.5	25726976	34663649280	7954	791	2024-08-24 23:21:17	3
6	0.2	25726976	34663649280	7954	791	2024-08-24 23:21:18	4

Function Entry/Exit logs

1	2024-08-24 23:21:17	- <START17049>
2	2024-08-24 23:21:18	- <END17049>
3	2024-08-24 23:21:18	- <START19555>
4	2024-08-24 23:21:28	- <END19555>
5	2024-08-24 23:21:28	- <START41611>
6	2024-08-24 23:21:38	- <END41611>
7	2024-08-24 23:21:38	- <START0006>

Comparison
Based on time


Creating Analytics
-> Total CO2
-> Total Energy Consumed
-> Elapsed time per module etc.

Energy Consumption: $E_{CPU} = \text{CPU Utilization}(\%) \times \text{CPU Power Consumption}(W) \times \text{Time}(h)$

$E_{Memory} = \left(\frac{RSS+VMS}{2} \right) \times \text{Memory Power Consumption per Byte}(W/\text{byte}) \times \text{Time}(h)$

Total Energy Consumption(E) = $E_{CPU} + E_{Memory}$

CO2 Emissions: $\text{CO2 Emissions} = \text{Total Energy Consumption}(E) \times \text{Emission Factor}$



Critical Code Blocks
Where Energy/CO2 >
Threshold



Optimized Code Suggestion Using
Code Llama


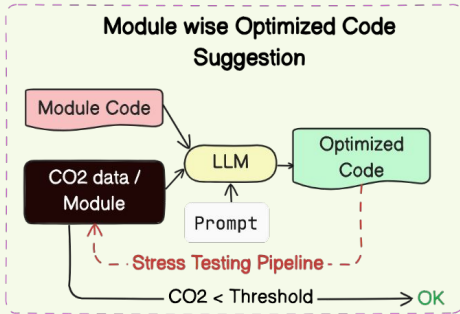
opentext™

Email / Notification
Service



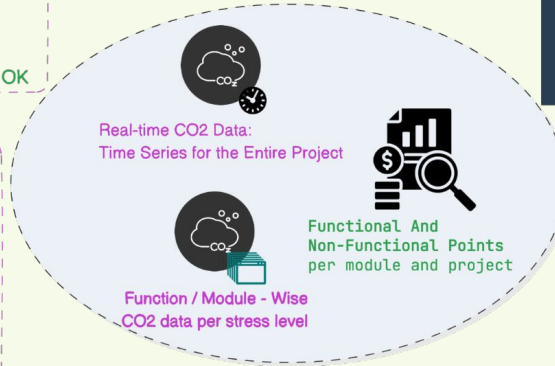

Module Wise parameter
visualization in dashboard

Data Utilisation & Processing (1/2)





opentext™

</> GREEN CODING;



**Email Notification in case of
Anomaly (OpenText)**



Data Utilisation & Processing

(2/2)

We mainly use three kinds of data :

1. Real-time CO2 Data: Time Series for the Entire Project
2. Function / Module - Wise CO2 data per stress level
3. Functional And Non-Functional Points per module and project

Here are some use cases we offer:

1. Module wise resource estimation
2. Optimized Code Suggestion
3. Anomaly Detection to prevent potential Faults
4. Scheduled Reports
5. Notification Alerts for high CO2 emission and energy use

TOOLS REQUIRED

Services

1. OpenText
2. green-coding

AI Frameworks

1. Python
2. LlamaIndex
3. codellama

Web Frameworks and databases

1. NextJS (For frontend and dashboards)
2. Flask (For backend API design)



USER EXPERIENCE

Welcome to CarbonSage

Your tool for tracking and reducing carbon emissions. Upload your project to get started.

Upload Your Project

Select a folder to upload. This will initialize your session and allow you to access the dashboard.

Choose Files no files selected

Upload

Files

```
../target_projects/proj_526c  
4558-43a7-adf6-  
2d9e15db073f/proj1/lib1.py  
../target_projects/proj_526c  
4558-43a7-adf6-  
2d9e15db073f/proj1/lib2.py  
../target_projects/proj_526  
4558-43a7-adf6-  
2d9e15db073f/proj1/main,
```

Dashboard

Session ID: 526ca59f-4558-43a7-adf6-2d9e15db073f

Upload Another Project

Process Project Folder

File: ../target_projects/proj_526ca59f-4558-43a7-adf6-2d9e15db073f/proj1/main.py

```
import logging
logging.basicConfig(filename='../function_logs/proj1.log', level=logging.INFO, format='%(%asctime)s - %(message)s',
datefmt='%Y-%m-%d %H:%M:%S')
from lib1 import *
from lib2 import *
import argparse

if __name__ == '__main__':
    # Current process ID
    print(f"Current process ID: {os.getpid()}")
    parser = argparse.ArgumentParser(description="Process some arguments.")
    parser.add_argument('--cpu_power', type=int, help='Power of CPU bound task')
    parser.add_argument('--memory_power', type=int, help='Power of memory bound task')
    parser.add_argument('--io_power', type=int, help='Power of I/O bound task')
    parser.add_argument('--power_of_10', type=int, help='Power of 10 for sorting list')

    args = parser.parse_args()

    if args.cpu_power:
        cpu_bound(args.cpu_power)

    print("Exiting...")
```

Parameters

Power of CPU bound task (cpu_power)

Value: 0

0

10

Power of memory bound task (memory_power)

Value: 0

0

10

Power of I/O bound task (io_power)

Value: 0

0

10

Power of 10 for sorting list (power_of_10)

Value: 0

0

10

Start Profiling

Profiling Status

No process to analyze.

Files

```
../target_projects/proj_70b8  
73d9-41cd-9dd6-  
02652e4754c8/lib1.py  
../target_projects/proj_70b8  
73d9-41cd-9dd6-  
02652e4754c8/lib2.py  
../target_projects/proj_70b  
73d9-41cd-9dd6-  
02652e4754c8/main.py
```

Dashboard

Session ID: 70b8828a-73d9-41cd-9dd6-02652e4754c8

Upload Another Project

Process Project Folder

File: ../target_projects/proj_70b8828a-73d9-41cd-9dd6-02652e4754c8/main.py

```
import logging
logging.basicConfig(filename='../function_logs/proj1.log', level=logging.INFO, format='%(%asctime)s - %(message)s',
datefmt='%Y-%m-%d %H:%M:%S')
from lib1 import *
from lib2 import *
import argparse

if __name__ == '__main__':
    # Current process ID
    print(f"Current process ID: {os.getpid()}")
    parser = argparse.ArgumentParser(description="Process some arguments.")
    parser.add_argument('--cpu_power', type=int, help='Power of CPU bound task')
    parser.add_argument('--memory_power', type=int, help='Power of memory bound task')
    parser.add_argument('--io_power', type=int, help='Power of I/O bound task')
    parser.add_argument('--power_of_10', type=int, help='Power of 10 for sorting list')

    args = parser.parse_args()

    if args.cpu_power:
        cpu_bound(args.cpu_power)
```

Analytics Completed

The results have been sent to
developerkabr@gmail.com.
Powered by OpenText API.

Close

USER EXPERIENCE

Function Profile Map

FUNCTION ID	CPU %	PAGEINS	PFAULTS	PHYSICAL MEMORY (RSS MB)	VIRTUAL MEMORY (VMS MB)	ELAPSED TIME (SECS)	ENERGY CONSUMPTION (KWH)	CO2 EMISSIONS (KG)
4247	0.00%	165	2079	22.66 MB	400848.14 MB	1.00 secs	0.00 kWh	0.000029 kg
8550	0.00%	165	2079	22.66 MB	400848.14 MB	1.00 secs	0.00 kWh	0.000029 kg
9412	0.00%	165	2079	22.66 MB	400848.14 MB	1.00 secs	0.00 kWh	0.000029 kg
9996	0.00%	165	2079	22.66 MB	400848.14 MB	1.00 secs	0.00 kWh	0.000029 kg
13634	0.00%	165	2079	22.66 MB	400848.14 MB	1.00 secs	0.00 kWh	0.000029 kg
13723	0.00%	165	2079	22.66 MB	400848.14 MB	1.00 secs	0.00 kWh	0.000029 kg
14331	0.00%	165	2079	22.66 MB	400848.14 MB	1.00 secs	0.00 kWh	0.000029 kg
17046	0.00%	165	2079	22.53 MB	400830.14 MB	1.00 secs	0.00 kWh	0.000030 kg

Analytics Dashboard

Data from Aug 25, 2024 22:56:33 to Aug 25, 2024 22:56:53

Average Values For Function Parameters

METRIC	VALUE
AVG CPU PERCENT	0.875000
AVG PAGEINS	165.000000
AVG PFAULTS	2077.409091
AVG RSS	23730734.545455
AVG VMS	420315986478.545471
CO2 EMISSIONS KG	0.000102
ELAPSED TIME SECS	3.500000
ENERGY CONSUMPTION KWH	0.000205

Code from ../target_projects/proj_70b8828a-73d9-41cd-9dd6-02652e4754c8/lib1.py

Optimize Code

Original Code:

```
# Memory Heavy Task

def memory_bound(p):
    logging.info('')
    print("Starting memory bound task")
    x = []
    for i in range(10**p):
        x.append(i)
        if i % 10**((p-1)) == 0:
            print(f"Memory bound task: {i}")
            time.sleep(1)
    print("Exiting memory bound task")
    logging.info('')
    return x

# I/O Heavy Task

def io_bound(p):
    logging.info('')
    # Generate a large file
    print("Starting I/O bound task")
    with open("large_file.txt", "w") as f:
        for i in range(10**p):
```

Optimized Code:

```
reduce the memory usage as we're not storing the entire range of
numbers in memory at once.

Here's the optimized code snippet:
...
from itertools import islice

def memory_bound(p):
    logging.info('')
    print("Starting memory bound task")
    x = []
    for i in islice(xrange(10**p), 1000):
        x.append(i)
        if i % 10**((p-1)) == 0:
            print(f"Memory bound task: {i}")
            time.sleep(1)
    print("Exiting memory bound task")
    logging.info('')
    return x
...
```

In this optimized version, we're using the 'xrange' function to generate the range of numbers and the 'islice' function to iterate over the range of numbers in chunks of 100

Close

pageins

From 25/8/2024, 10:56:33 PM to 25/8/2024, 10:56:53 PM



Min: 165.00

Max: 165.00

Average: 165.00

pfaults

From 25/8/2024, 10:56:33 PM to 25/8/2024, 10:56:53 PM



Min: 2071.00

Max: 2079.00

Average: 2076.71

Rss

From 25/8/2024, 10:56:33 PM to 25/8/2024, 10:56:53 PM




Vms

From 25/8/2024, 10:56:33 PM to 25/8/2024, 10:56:53 PM



IMPACT

30%




Cost Savings

Reducing energy consumption can save businesses millions annually (Energy Star).


Environmental Benefits

Green technologies can reduce carbon emissions (McKinsey & Company).



15%

8%




Data Driven Innovation

Leveraging data analytics can achieve higher operational efficiency (Deloitte).

Competitive Advantage

Global consumers are willing to pay more for sustainable brands (Nielsen).



66%


FUTURE WORKS

Integration with Cloud Providers

Develop plugins and APIs for seamless integration with major cloud platforms (AWS, Azure, GCP).

Data Security and Privacy

Ensure compliance with data protection regulations (e.g., GDPR, CCPA).



Regulatory Compliance

Help users comply with global and local emission regulations. Provide regular updates on regulatory changes and their impact on emissions tracking.

Thank You

